



Machine Learning

# Application example: Photo OCR

---

## Problem description and pipeline

# The Photo OCR problem

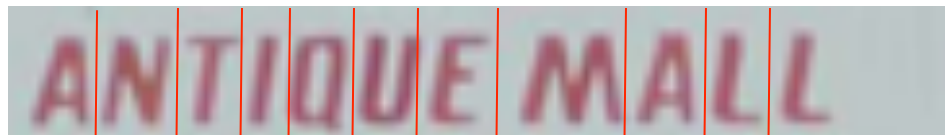


# Photo OCR pipeline

→ 1. Text detection



→ 2. Character segmentation

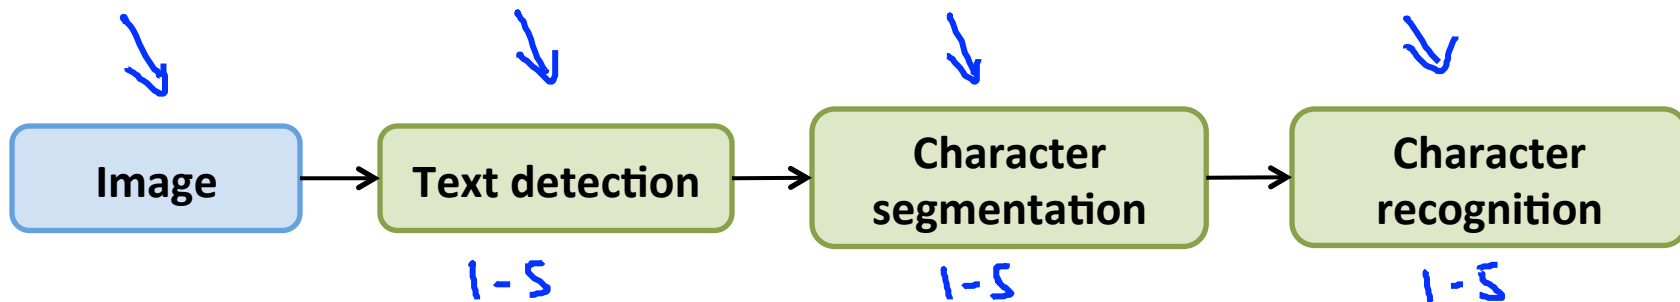


→ 3. Character classification



~~Cleaning~~ → ~~Cleaning~~

# Photo OCR pipeline





Machine Learning

Application example:  
Photo OCR


---

Sliding windows

## Text detection



## Pedestrian detection



# Supervised learning for pedestrian detection

$x$  = pixels in 82x36 image patches

1,000  
10,000  
...



Positive examples ( $y = 1$ )



Negative examples ( $y = 0$ )



# Sliding window detection

step-size / stride






# Sliding window detection



# Sliding window detection



# Sliding window detection



# Text detection



# Text detection



Positive examples ( $y = 1$ )

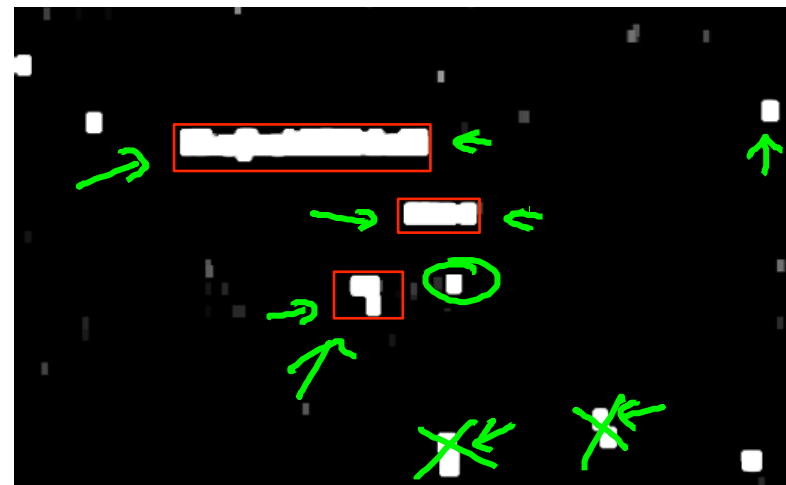




Negative examples ( $y = 0$ )



# Text detection



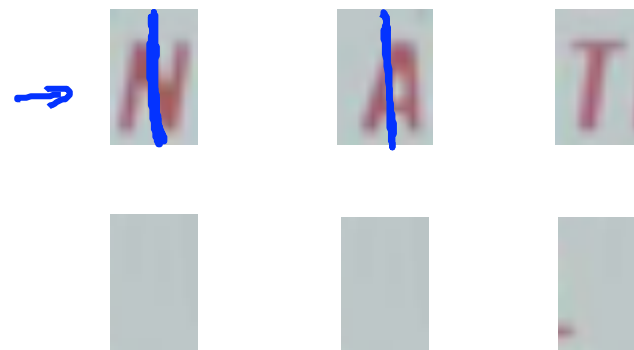
"expansion"



# 1D Sliding window for character segmentation



Positive examples ( $y = 1$ )



Negative examples ( $y = 0$ )

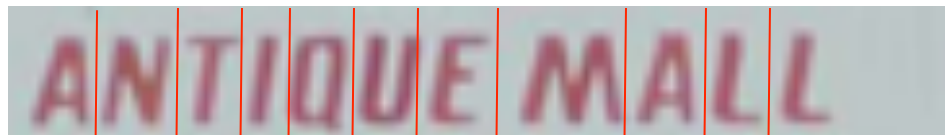


# Photo OCR pipeline

→ 1. Text detection



→ 2. Character segmentation



→ 3. Character classification





Machine Learning

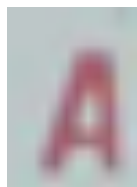


# Application example: Photo OCR

---

## Getting lots of data: Artificial data synthesis

# Character recognition



→ A



→ N



→ T



→ I



→ Q



→ A

# Artificial data synthesis for photo OCR



Real data



Abcdefg

Abcdefg

Abcdefg

Abcdefg

Abcdefg

# Artificial data synthesis for photo OCR





Real data



Synthetic data

# Synthesizing data by introducing distortions



# Synthesizing data by introducing distortions: Speech recognition



Original audio: 



Audio on bad cellphone connection



Noisy background: Crowd

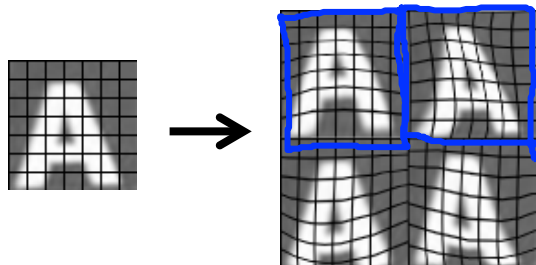


Noisy background: Machinery




# Synthesizing data by introducing distortions

- Distortion introduced should be representation of the type of noise/distortions in the test set.



- Audio:  
Background noise,  
bad cellphone connection

- Usually does not help to add purely random/meaningless noise to your data.



- $x_i$  = intensity (brightness) of pixel  $i$
- $x_i \leftarrow x_i + \text{random noise}$

## Discussion on getting more data

1. Make sure you have a low bias classifier before expending the effort. (Plot learning curves). E.g. keep increasing the number of features/number of hidden units in neural network until you have a low bias classifier.

2. “How much work would it be to get 10x as much data as we currently have?”

- Artificial data synthesis
- Collect/label it yourself
- “Crowd source” (E.g. Amazon Mechanical Turk)

→ #hours?  $n = 1,000$   
→ 10 secs/example  
 $n = 10,000$  ←

## Discussion on getting more data

1. Make sure you have a low bias classifier before expending the effort. (Plot learning curves). E.g. keep increasing the number of features/number of hidden units in neural network until you have a low bias classifier.
2. “How much work would it be to get 10x as much data as we currently have?”
  - Artificial data synthesis
  - Collect/label it yourself
  - “Crowd source” (E.g. Amazon Mechanical Turk)



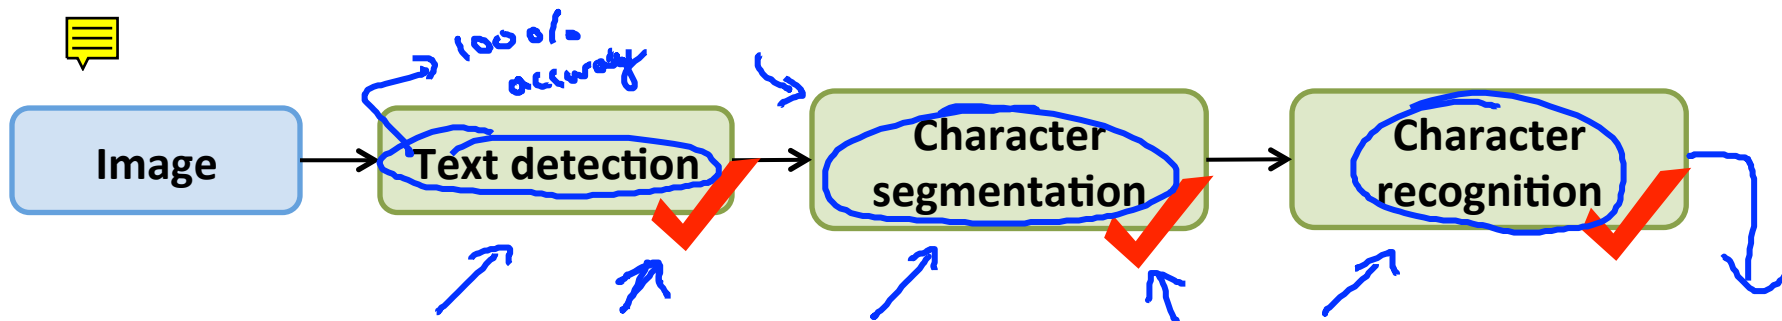
Machine Learning

# Application example: Photo OCR

---

Ceiling analysis: What  
part of the pipeline to  
work on next

# Estimating the errors due to each component (ceiling analysis)



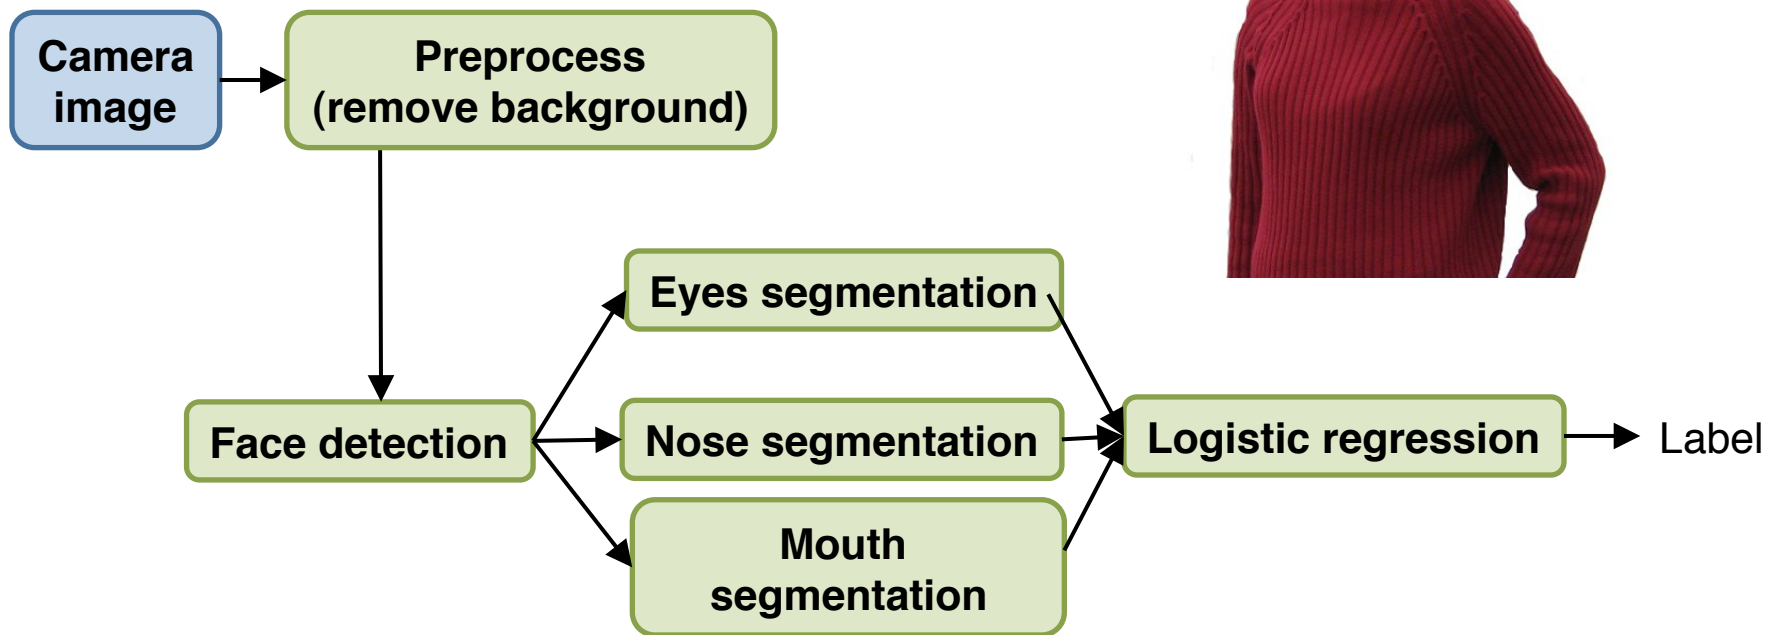
What part of the pipeline should you spend the most time trying to improve?

Component	Accuracy
Overall system	72%
→ Text detection	89%
Character segmentation	90%
Character recognition	100%

Handwritten annotations on the table include blue arrows pointing to the 'Accuracy' column and the 'Text detection' row, and blue arrows pointing to the 'Overall system', 'Text detection', 'Character segmentation', and 'Character recognition' rows. A blue arrow also points to the '100%' value in the 'Character recognition' row. A blue arrow points to the '72%' value in the 'Overall system' row. A blue arrow points to the '89%' value in the 'Text detection' row. A blue arrow points to the '90%' value in the 'Character segmentation' row. A blue arrow points to the '100%' value in the 'Character recognition' row. A blue arrow points to the '17%' value in the 'Overall system' row. A blue arrow points to the '1%' value in the 'Text detection' row. A blue arrow points to the '10%' value in the 'Character recognition' row.

## Another ceiling analysis example

Face recognition from images  
(Artificial example)



## Another ceiling analysis example

